

因为没有使用Maven管理项目，就没照着老师给的pdf中要求的提交格式交了

output文件夹中是输出文件，code文件夹中是代码文件（包括.java和.jar）

## 0 运行流程介绍

### 0.1 数据集准备

先查看所有容器 `sudo docker ps`（获得容器id）

```
chenmiao@chenmiao-virtual-machine:~$ sudo docker ps
CONTAINER ID   IMAGE          COMMAND                  CREATED        STATUS        PORTS
cb46a3dead24   newuhadoop    "/bin/bash"             8 days ago    Up 42 seconds
2758f3a05bf8   newuhadoop    "/bin/bash"             8 days ago    Up 42 seconds
e45c02dda942   newuhadoop    "/bin/bash"             8 days ago    Up 43 seconds
1d8cd276f450   newuhadoop    "/bin/bash"             8 days ago    Up 43 seconds
63f0da8c80fc   newuhadoop    "/bin/bash"             8 days ago    Up 44 seconds
```

将数据集传输到docker容器（h01）

`sudo docker cp <需要传输的文件路径> <容器id>:<传输文件在容器中的存放位置>`

```
chenmiao@chenmiao-virtual-machine:~$ sudo docker cp ~/financial_big_data/homework5/analyst_
ratings.csv 63f0da8c80fc:/tmp
Successfully copied 52.5MB to 63f0da8c80fc:/tmp
chenmiao@chenmiao-virtual-machine:~$ sudo docker cp ~/financial_big_data/homework5/stop-wor
d-list.txt 63f0da8c80fc:/tmp
Successfully copied 4.1kB to 63f0da8c80fc:/tmp
```

将数据集放到hdfs上（需要先在hdfs新建一个input文件夹）然后在usr/local/hadoop目录下执行：

`./bin/hdfs dfs -put <数据集在docker中存放的位置> /input/`

可以通过 `./bin/hdfs dfs -ls /input` 查看是否放好了

```
root@h01:/usr/local/hadoop# ./bin/hdfs dfs -put /tmp/analyst_ratings.csv /input/
root@h01:/usr/local/hadoop# ./bin/hdfs dfs -ls /input
Found 1 items
-rw-r--r--  2 root supergroup  52462980 2024-10-22 13:53 /input/analyst_rating
s.csv
root@h01:/usr/local/hadoop# ./bin/hdfs dfs -put /tmp/stop-word-list.txt /input/
root@h01:/usr/local/hadoop# ./bin/hdfs dfs -ls /input
Found 2 items
-rw-r--r--  2 root supergroup  52462980 2024-10-22 13:53 /input/analyst_rating
s.csv
-rw-r--r--  2 root supergroup      2231 2024-10-22 13:55 /input/stop-word-list
.txt
root@h01:/usr/local/hadoop#
```

### 0.2 运行代码准备

同理，先把java代码文件放到docker（这里我直接放到了hadoop文件夹下）

```
chenmiao@chenmiao-virtual-machine:~$ sudo docker cp ~/financial_big_data/homework5/StockCou
nt.java 63f0da8c80fc:/usr/local/hadoop
[sudo] password for chenmiao:
Successfully copied 5.63kB to 63f0da8c80fc:/usr/local/hadoop
```

然后生成相应的.class文件

```
1 javac -classpath `hadoop classpath` StockCount.java
```

再生成相应的.jar

```
1 | jar -cvf stockcount.jar StockCount*.class
```

### 0.3 MapReduce运行

```
1 # 把结果保存到/output/<新建文件夹>
2 ./bin/hadoop jar stockcount.jar StockCount /input/analyst_ratings.csv
  /output/<文件夹名>
3 # 查看
4 ./bin/hadoop fs -ls /output/<文件夹名>
5 # 打印输出
6 ./bin/hadoop fs -cat <输出文件的路径>
```

## 0.4 把docker中的文件传回

先在容器中运行以下命令，将文件从HDFS下载到容器的本地目录（代码文件不需要这步，输出文件需要）

```
./bin/hdfs dfs -get /output/stock_count/part-r-00000 <需要在docker中存放的路径>
```

再将该文件从 Docker 容器传输到本机。

```
docker cp <容器id>:<文件在docker中存放的位置> <文件在本机需要存储的位置>
```

## 0.5 遇到错误

生成.java文件相应的.class文件时，遇到多次报错。

**错误1:** StockCount.java:19: error: unmappable character for encoding ASCII

```
// ?????????????????????? ^
```

原因：代码文件中包含非ASCII编码的字符，编译器默认使用 ASCII 编码，无法识别这些字符。

解决方法：指定编码方式

```
1 javac -encoding UTF-8 -classpath `hadoop classpath` StockCount.java
```

**错误2:** StockCount.java:2: error: package org.apache.hadoop.conf does not exist

```
import org.apache.hadoop.conf.Configuration; ^
```

原因：找不到 Hadoop 的核心类库。

解决方法:

直接输入 `/usr/local/hadoop/bin/hadoop classpath`，会得到输出hadoop的classpath，然后直接把它粘贴到命令行里即可，最终完整可用的命令如下

```
1 javac -encoding UTF-8 -classpath
   "/usr/local/hadoop/etc/hadoop:/usr/local/hadoop/share/hadoop/common/lib/*:/usr
   /local/hadoop/share/hadoop/common/*:/usr/local/hadoop/share/hadoop/hdfs:/usr/l
   ocal/hadoop/share/hadoop/hdfs/lib/*:/usr/local/hadoop/share/hadoop/hdfs/*:/usr
   /local/hadoop/share/hadoop/mapreduce/*:/usr/local/hadoop/share/hadoop/yarn:/us
   r/local/hadoop/share/hadoop/yarn/lib/*:/usr/local/hadoop/share/hadoop/yarn/*"
   StockCount.java
```

# 1 StockCount

要求：在HDFS上加载上市公司热点新闻标题数据集（analyst\_ratings.csv），该数据集收集了部分上市公司的热点财经新闻标题。编写MapReduce程序完成以下任务：统计数据集上市公司股票代码（“stock”列）的出现次数，按出现次数从大到小输出，输出格式为“<排名>: <股票代码>, <次数>”。

## 1.1 设计思路

最直观的想法是先将需要的股票代码（stock列）提取出来，维护一个含股票代码和出现次数两个字段的列表，对重复出现的股票代码进行次数累加。最后做排序和格式的修改。

- Mapper

定义一个StockMapper类，逐行处理数据，使用split方法切片，提取第四个字段stock，然后将每个股票代码视为出现了1次，输出键值对 <股票代码, 1>。

- Reducer

定义一个StockReducer类，接收来自Mapper的输出，累加相同股票代码的出现次数，得到每个股票代码的总出现次数；并维护一个列表存储每个股票代码及其出现次数。

对列表按照股票代码出现次数进行降序排序，遍历排序后的列表，为每个股票代码生成一个排名，并按照指定的格式输出结果。

- main

设置MapReduce作业的配置：配置Mapper和Reducer类；指定输入、输出文件路径（args[0]和args[1]）；启动作业并等待作业完成。

## 1.2 程序运行结果

- 终端输出结果

运行MapReduce结果

```
root@h01:/usr/local/hadoop# ./bin/hadoop jar stockcount.jar StockCount /input/analyst_ratings.csv /output/stock_count
2024-10-22 15:38:13,787 INFO client.DefaultNoHARMFailoverProxyProvider: Connecting to ResourceManager at h01/172.18.0.2:8032
2024-10-22 15:38:14,501 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
2024-10-22 15:38:14,539 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/root/.staging/job_1729604884725_0002
2024-10-22 15:38:15,021 INFO input.FileInputFormat: Total input files to process : 1
2024-10-22 15:38:15,165 INFO mapreduce.JobSubmitter: number of splits:1
2024-10-22 15:38:15,364 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1729604884725_0002
2024-10-22 15:38:15,364 INFO mapreduce.JobSubmitter: Executing with tokens: []
2024-10-22 15:38:15,679 INFO conf.Configuration: resource-types.xml not found
2024-10-22 15:38:15,680 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2024-10-22 15:38:15,837 INFO impl.YarnClientImpl: Submitted application application_1729604884725_0002
2024-10-22 15:38:15,907 INFO mapreduce.Job: The url to track the job: http://h01
```

```
root@h01: /usr/local/hadoop
Merged Map outputs=1
GC time elapsed (ms)=709
CPU time spent (ms)=7090
Physical memory (bytes) snapshot=789815296
Virtual memory (bytes) snapshot=5194358784
Total committed heap usage (bytes)=756547584
Peak Map Physical memory (bytes)=545624064
Peak Map Virtual memory (bytes)=2593247232
Peak Reduce Physical memory (bytes)=244191232
Peak Reduce Virtual memory (bytes)=2601111552

Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=52462980
File Output Format Counters
  Bytes Written=86831
root@h01:/usr/local/hadoop# ./bin/hadoop fs -ls /output/stock_count
```

打印输出文件结果

```
root@h01:/usr/local/hadoop# ./bin/hadoop fs -ls /output/stock_count
Found 2 items
-rw-r--r--  2 root supergroup      0 2024-10-22 15:38 /output/stock_count/_SUCCESS
-rw-r--r--  2 root supergroup 86831 2024-10-22 15:38 /output/stock_count/part-r-00000
root@h01:/usr/local/hadoop# ./bin/hadoop fs -cat /output/stock_count/part-r-00000
1: MS, 726
2: MRK, 704
3: QQQ, 693
4: BABA, 689
5: EWU, 681
6: GILD, 663
7: JNJ, 663
8: MU, 659
9: NVDA, 655
10: VZ, 648
11: KO, 643
12: QCOM, 636
13: M, 635
14: NFLX, 635
15: EBAY, 621
16: DAL, 605
```

- WEB页面截图

## Browse Directory

Show 25 entries

Search:

<input type="checkbox"/>	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	<input type="checkbox"/>
<input type="checkbox"/>	drwxr-xr-x	root	supergroup	0 B	Oct 22 21:55	0	0 B	input	<input type="checkbox"/>
<input type="checkbox"/>	drwxr-xr-x	root	supergroup	0 B	Oct 22 23:38	0	0 B	output	<input type="checkbox"/>
<input type="checkbox"/>	drwxr-xr-x	root	supergroup	0 B	Oct 14 18:17	0	0 B	tmp	<input type="checkbox"/>
<input type="checkbox"/>	drwxr-xr-x	root	supergroup	0 B	Oct 14 18:09	0	0 B	user	<input type="checkbox"/>

Showing 1 to 4 of 4 entries

Previous

1

Next

Hadoop, 2024.

## Browse Directory

Show 25 entries

Search:

<input type="checkbox"/>	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	<input type="checkbox"/>
<input type="checkbox"/>	-rw-r--r--	root	supergroup	0 B	Oct 22 23:38	2	128 MB	_SUCCESS	<input type="checkbox"/>
<input type="checkbox"/>	-rw-r--r--	root	supergroup	84.8 KB	Oct 22 23:38	2	128 MB	part-r-00000	<input type="checkbox"/>

Showing 1 to 2 of 2 entries

Previous

1

Next

Hadoop, 2024.

## 1.3 程序分析

进一步对性能、扩展性等方面存在的不足和可能改进之处进行分析。

### 1.3.1 Mapper输出量

**不足：**对每一行（股票代码不为空的）数据都输出 `<股票代码, 1>`，导致Mapper输出量太大。

**改进分析：**先在Mapper对股票代码的计数进行局部聚合，减少Mapper和Reducer之间的数据传输量。

### 1.3.2 Reducer运行内存

**不足：**使用List来存储所有股票代码和它们的出现次数，并对整个List进行排序，这对于大型数据集来说会占用大量内存，可能导致崩溃。

**改进分析：**可以将数据分成多个部分进行排序，然后通过归并排序的方法将结果合并，降低内存占用。

## 2 WordFrequency

要求：在HDFS上加载上市公司热点新闻标题数据集（analyst\_ratings.csv），该数据集收集了部分上市公司的热点财经新闻标题。统计数据集热点新闻标题（“headline”列）中出现的前100个高频单词，按出现次数从大到小输出。要求忽略大小写，忽略标点符号，忽略停词（stop-word-list.txt）。输出格式为“<排名>: <单词>, <次数>”。

### 2.1 设计思路

设计的思路与前一个任务基本一致，但多了几个处理：停词表的加载、对切片后字段的处理，

- Mapper

将停词表加载到HashSet中，可以在后续操作快速查询是否为停词。

逐行处理数据，使用split方法切片，提取第二个字段headline；去除标点符号，将文本转换为小写，并按空格分割成单词；检查每个单词是否为停词或空词，如果不是，则将其输出为键值对 `<单词, 1>`。

- Reducer

接收来自Mapper的输出，累加相同单词的出现次数，得到每个单词的总出现次数；并维护一个map存储每个单词及其出现次数。

cleanup方法在Reducer任务结束后将map转换为list，按照单词出现次数进行降序排序，输出出现次数最多的前100个单词，格式为“<排名>: <单词>, <次数>”。

- main

设置MapReduce作业的配置：配置Mapper和Reducer类；指定输入、输出文件路径（args[0]和args[1]）；启动作业并等待作业完成。

### 2.2 程序运行结果

- 终端输出结果

运行MapReduce结果

```
root@h01: /usr/local/hadoop
root@h01:/usr/local/hadoop# ./bin/hadoop jar wordfrequency.jar WordFrequency /input/analyst_ratings.csv /output/word_frequency /input/stop-word-list.txt
2024-10-22 18:25:04,818 INFO client.DefaultNoHARMFailoverProxyProvider: Connecting to ResourceManager at h01/172.18.0.2:8032
2024-10-22 18:25:05,450 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
2024-10-22 18:25:05,486 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/root/.staging/job_1729604884725_0004
2024-10-22 18:25:05,956 INFO input.FileInputFormat: Total input files to process : 1
2024-10-22 18:25:06,090 INFO mapreduce.JobSubmitter: number of splits:1
2024-10-22 18:25:06,274 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1729604884725_0004
2024-10-22 18:25:06,274 INFO mapreduce.JobSubmitter: Executing with tokens: []
2024-10-22 18:25:06,582 INFO conf.Configuration: resource-types.xml not found
2024-10-22 18:25:06,583 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2024-10-22 18:25:06,717 INFO impl.YarnClientImpl: Submitted application application_1729604884725_0004
2024-10-22 18:25:06,808 INFO mapreduce.Job: The url to track the job: http://h01:8088/proxy/application_1729604884725_0004/
2024-10-22 18:25:06,809 INFO mapreduce.Job: Running job: job_1729604884725_0004
2024-10-22 18:25:15,072 INFO mapreduce.Job: Job job_1729604884725_0004 running in uber mode : false
2024-10-22 18:25:15,073 INFO mapreduce.Job:  map 0% reduce 0%
2024-10-22 18:25:29,357 INFO mapreduce.Job:  map 100% reduce 0%
2024-10-22 18:25:38,468 INFO mapreduce.Job:  map 100% reduce 100%
```

```
Reduce output records=100
Spilled Records=4267858
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=568
CPU time spent (ms)=14130
Physical memory (bytes) snapshot=834555904
Virtual memory (bytes) snapshot=5201440768
Total committed heap usage (bytes)=783286272
Peak Map Physical memory (bytes)=590536704
Peak Map Virtual memory (bytes)=2594750464
Peak Reduce Physical memory (bytes)=244019200
Peak Reduce Virtual memory (bytes)=2606690304

Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=52462980

File Output Format Counters
  Bytes Written=1799
```

打印输出文件结果



```
root@h01:/usr/local/hadoop# ./bin/hadoop fs -ls /output/word_frequency/
Found 2 items
-rw-r--r--    2 root supergroup          0 2024-10-22 18:25 /output/word_frequency/_SUCCESS
-rw-r--r--    2 root supergroup      1799 2024-10-22 18:25 /output/word_frequency/part-r-000000
root@h01:/usr/local/hadoop# ./bin/hadoop fs -cat /output/word_frequency/part-r-000000
1: stocks, 37669
2: shares, 26843
3: q, 25950
4: vs, 24905
5: m, 24538
6: update, 23804
7: market, 22458
8: est, 20181
9: reports, 19300
10: eps, 18050
11: session, 14337
12: week, 14202
88: gathers, 3354
89: services, 3331
90: volume, 3264
91: ahead, 3201
92: midafternoon, 3190
93: etf, 3182
94: report, 3175
95: futures, 3173
96: wednesdays, 3163
97: thursdays, 3096
98: coverage, 3077
99: amid, 3049
100: ceo, 3007
```

- WEB页面截图

Browse Directory

/

Go!

Show

25

entries

Search:

	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	
<input type="checkbox"/>	drwxr-xr-x	root	supergroup	0 B	Oct 22 21:55	0	0 B	input	
<input type="checkbox"/>	drwxr-xr-x	root	supergroup	0 B	Oct 23 02:25	0	0 B	output	
<input type="checkbox"/>	drwxr-xr-x	root	supergroup	0 B	Oct 14 18:17	0	0 B	tmp	
<input type="checkbox"/>	drwxr-xr-x	root	supergroup	0 B	Oct 14 18:09	0	0 B	user	

Showing 1 to 4 of 4 entries

Previous

1

Next

/output/word\_frequency

Go!

Show

25

entries

Search:

	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	
<input type="checkbox"/>	-rw-r--r--	root	supergroup	0 B	Oct 23 02:25	2	128 MB	_SUCCESS	
<input type="checkbox"/>	-rw-r--r--	root	supergroup	1.76 KB	Oct 23 02:25	2	128 MB	part-r-00000	

Showing 1 to 2 of 2 entries

Previous

1

Next

Hadoop, 2024.



## 2.3 程序分析

---

进一步对性能、扩展性等方面存在的不足和可能改进之处进行分析。

### 2.3.1 Reducer内存

**不足：**WordReducer将所有的单词计数存储在一个HashMap中。如果输入数据量非常大，单词的种类和数量超过Reducer的内存限制，会导致内存溢出问题。

**改进分析：**先在Mapper对单词的计数进行局部聚合，减少Mapper和Reducer之间的数据传输量，以此减少Reducer阶段需要处理的数据量，从而减轻内存压力。

### 2.3.2 排序优化

**不足：**在Reducer的cleanup方法中，对所有单词进行排序，性能开销较大，尤其当单词总量较大时，排序时间会显著增长。

**改进分析：**可以先找到出现次数第100大的元素，然后剔除比该元素出现次数少的单词，对最高频100个元素进行排序，减少了不必要的全局排序开销。